

Week 13 Report

Kieran Grant - 2357351G

What I've done this week

- Submitted final draft of interim report.
- Implemented data generation and preprocessing pipeline. Seems to all be working well!
 - Pipeline randomly samples a song from the MusDB dataset, takes a 5 second sample of the song and applies two different effect settings to the audio.
 - Some augmentation is also applied to the sample before DAFX processing using built-in SoX effects - just some very basic pitch shifting and time stretching to increase the variability of the training set.
 - There are checks that the output audio is not all silent, and that the transformation is sensible (using cosine similarity between original and effected MFCCs - manually tuned this by listening to some of the generated data).
 - The data generation is pretty fast as well so shouldn't be a bottleneck during training.
- Added Audio Commons feature extractor to preprocessing pipeline. A few issues with this:
 - Feature extraction is slow, even on these 5 second samples. To extract all 8 features for a single audio sample takes around 3 seconds. Considering that I was going to use a training set size of 10,000 (and each set has 2 audio samples) - this would mean an epoch would take at least 16.67 hours.
 - The extraction models themselves are not the most robust and frequently throw errors (~15% of cases for most features). I've wrapped each in a try/catch block for now which returns a default value for that feature if an error occurs.
 - To reduce the amount of time needed for feature extraction I did some analysis on the correlations and variance of each of the features to select a subset to use for training. I removed any instance where a default value was returned and also removed the 'reverb' feature as this had very little variance across all the samples.
 - Warmth had a large (negative) correlation with brightness and positive correlation with boominess so I decided to remove this feature. It was also the most flaky - returning a default value in ~60% of samples.
 - The final (absolute) correlation matrix can be found in Figure 1. I decided to use only brightness and depth for now as features to be extracted for training. This brings the feature extraction time down to ~0.5 seconds for a 5 second piece of audio - still not great, but better than extracting all features!
- Implemented a 'first draft' of training pipeline for both the β -VAE and the end-to-end system.
 - The good news is the VAE is able to be trained on the GPU as it doesn't use SPSA.
 - * I performed some very short training runs to check that the VAE was learning during training.
 - * A summary of one of these runs can be found here:
 - [Link to W&B Report](#)
 - * Encouraging to see that the model seems to be learning quite quickly even on this very small dataset.
 - * I used a slightly larger latent space (64-dimensions) for the VAE which can be further mapped down in the end-to-end system.
 - I then used the learned weights of the VAE encoder section in the end-to-end system.
 - * Seems to load in the weights correctly and runs training end-to-end including SPSA for the DAFX.
 - * I haven't implemented the 'proper' loss function (just using MSE for now) so not worth sharing the training logs.
 - * Still having the same issue as last week with running model training on the GPU with SPSA. However, training on the CPU seems reasonably quick so probably not an issue.

Questions

I would be interested to hear your thoughts on the following:

- For the VAE training, I need to use a DAFX for audio generation and I am thinking about the best way to do this:
 - Could use either no DAFX (just pitch and time-stretch augmentation) or something very transparent (compressor with only one or two parameters being changed slightly).
 - Use different DAFX at each epoch - not sure if this would make training unstable? Also not sure if the model would learn anything more than just seeing lots of examples of audio.
- I still need to decide on exactly how to map from the learned VAE latent space down to a low-dimensional latent space then to parameter settings. A possible approach is shown in Figure 2. Though for making sure the low-dimensional latent space is sensible I will only be able to use the KL-divergence and not reconstruction loss since the output of the decoder is the parameter settings.

Plan for next semester

- Decide on audio generation for training VAE.
- Train the VAE and perform some analysis on the learned latent space - tune the latent space dimension.
- Use learned weights for end-to-end model.
- Decide on mapping from learnt latent space to parameter settings.
- Implement sensible loss function for both the audio and the low-dimensional latent space.
- Train end-to-end model for a few different effects.

Current state of project

- Interim report complete.
- Made good implementation progress.
- Still a few decisions to be made around model training and mapping learned latent space to DAFX parameters for each new DAFX.




Figure 1: Audio Common feature correlation matrix (3715 samples across 9 different effects including compressor, distortion, reverb, delay and modulation).

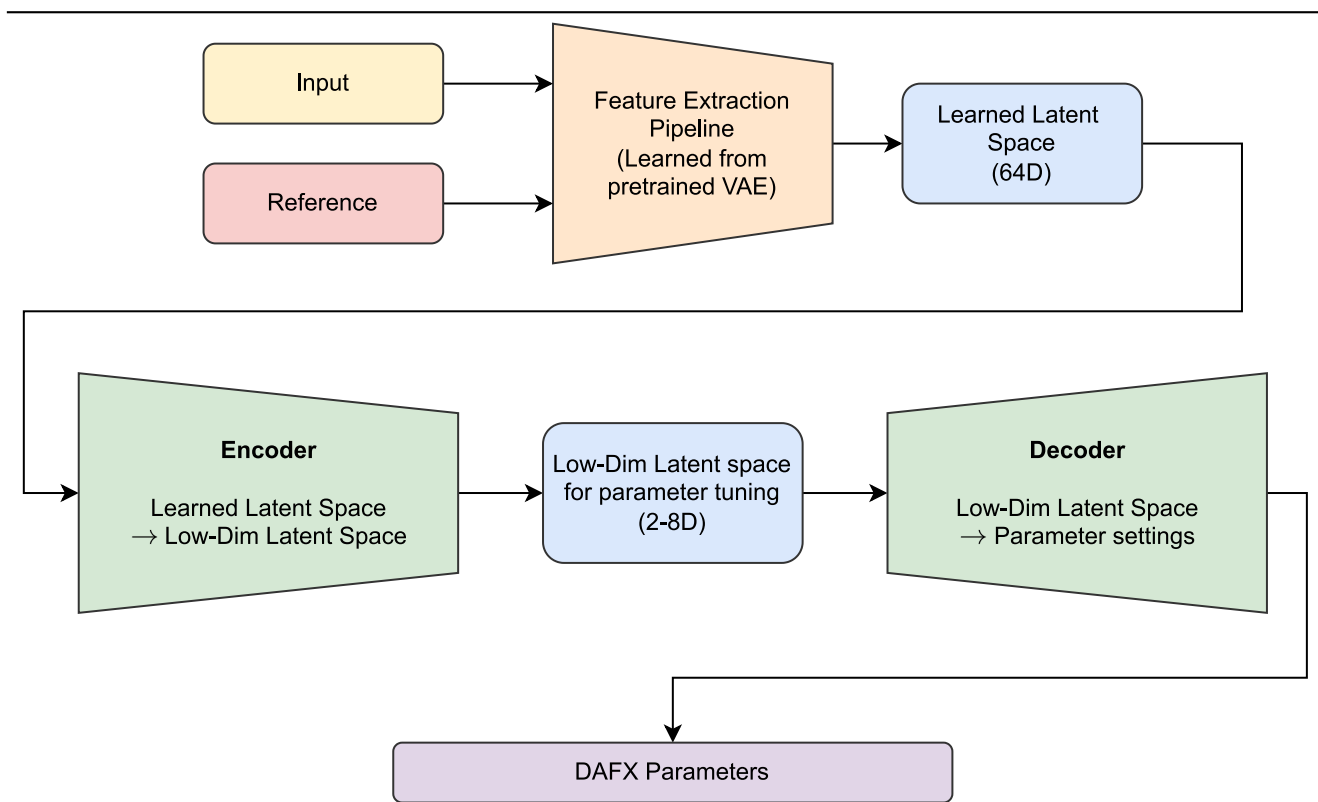


Figure 2: Rough idea of end-to-end training after VAE pre-training.
