

Week 5 Report

Kieran Grant - 2357351G

What I've done this week

- Updated Spectrogram-VAE model to take joint input-target audio pairs, stacked across channels.
 - Implemented 2-channel version for Spectrogram without phase and 4-channel version for Spectrogram with phase.
 - Had some issues when training with 4-channel version, decided to use 2-channel version for now.
 - Updated loss to use sum rather than batch-mean. Makes tracking training less interpretable, but model seems to train better.
 - Trained Joint-Spectrogram VAE on VTCK dataset with Delay and Overdrive DAFX.
 - * Train logs can be found here: <https://api.wandb.ai/links/kieran-grant/7iluvtrs>
 - * Training quite unstable at beginning - however, stabilised after ~50 epochs.
 - * I'm still not 100% sure why old loss was causing issues with KLD-loss. Still used a very small β value (0.0001) for Spectrogram-VAE training.
 - Reconstructions (**Figures 2** and **3** below) and latent space clustering (**Figure 1**) look promising.
 - Performed interpolation for audio and DAFX with learned model, results look sensible (**Figures 4** and **5**).
- Started implementation of simple end-to-end system, without bottleneck.
 - i.e. mapping joint embedding from Spectrogram VAE to DAFX parameters directly with a simple linear network.
 - Loss uses combination of Multiresolution STFT loss and L1 loss directly on audio as well as tracking DAFX parameter settings.
 - Created callback to store input/reference/prediction audio on W&B.
 - Ran short training session, had some issues with predicting correct settings with dummy values - will look more into this next week.

Questions

- Now that I have a simple trained Spectrogram-VAE - I'm wondering about the best method to map this down to a low dimensional (2-8D) latent space that can be 'played with' to find similar parameter settings. Would UMAP work, or should I have another bottleneck in a linear network and use similar sampling/KLD as in the autoencoder before mapping to parameter settings?

Plan for next week

- Debug end-to-end system.
- Implement latent bottleneck for controller network.

Current state of project

- Spectrogram VAE looks like it is performing a little better, should hopefully unblock me for the moment so I can work on the controller network.

Joint-Spectrogram VAE Visualisations




Figure 1: Learned latent space of Joint-Spectrogram VAE (2,000 samples per DAFX).

mda Overdrive

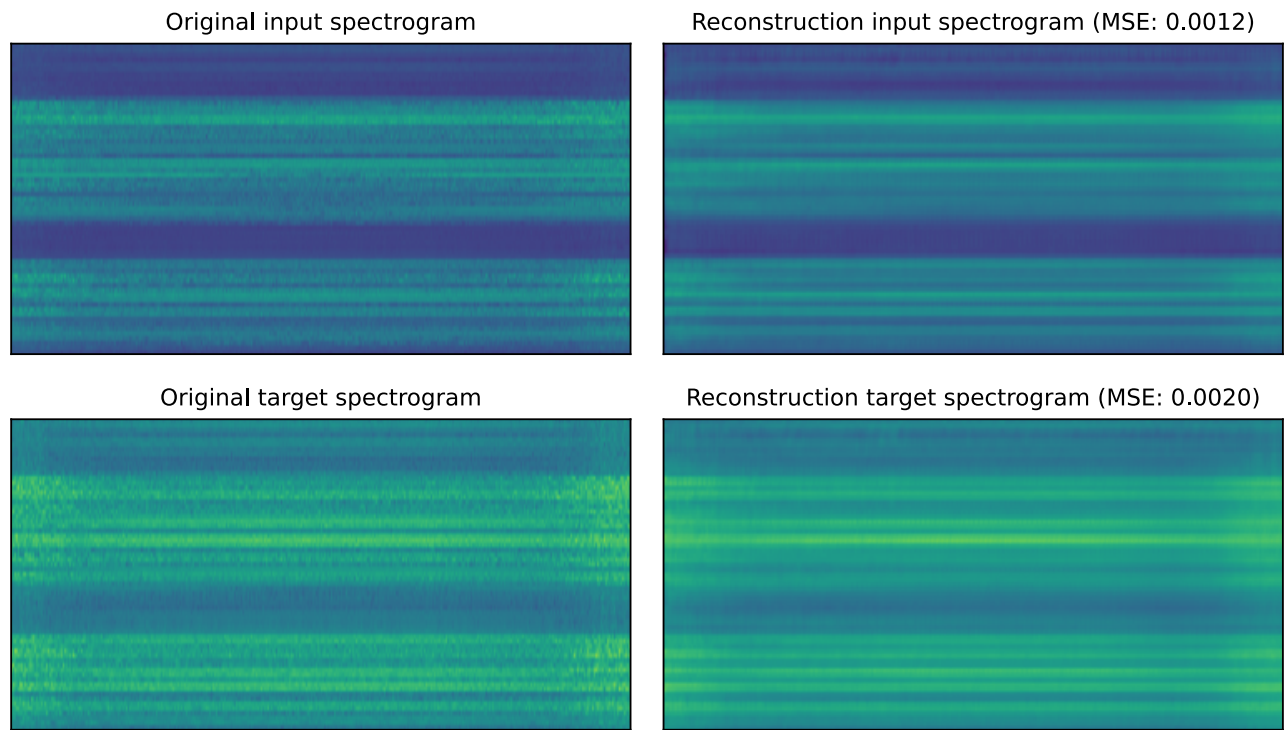


Figure 2: Reconstruction of Overdrive DAFX.

mda Delay

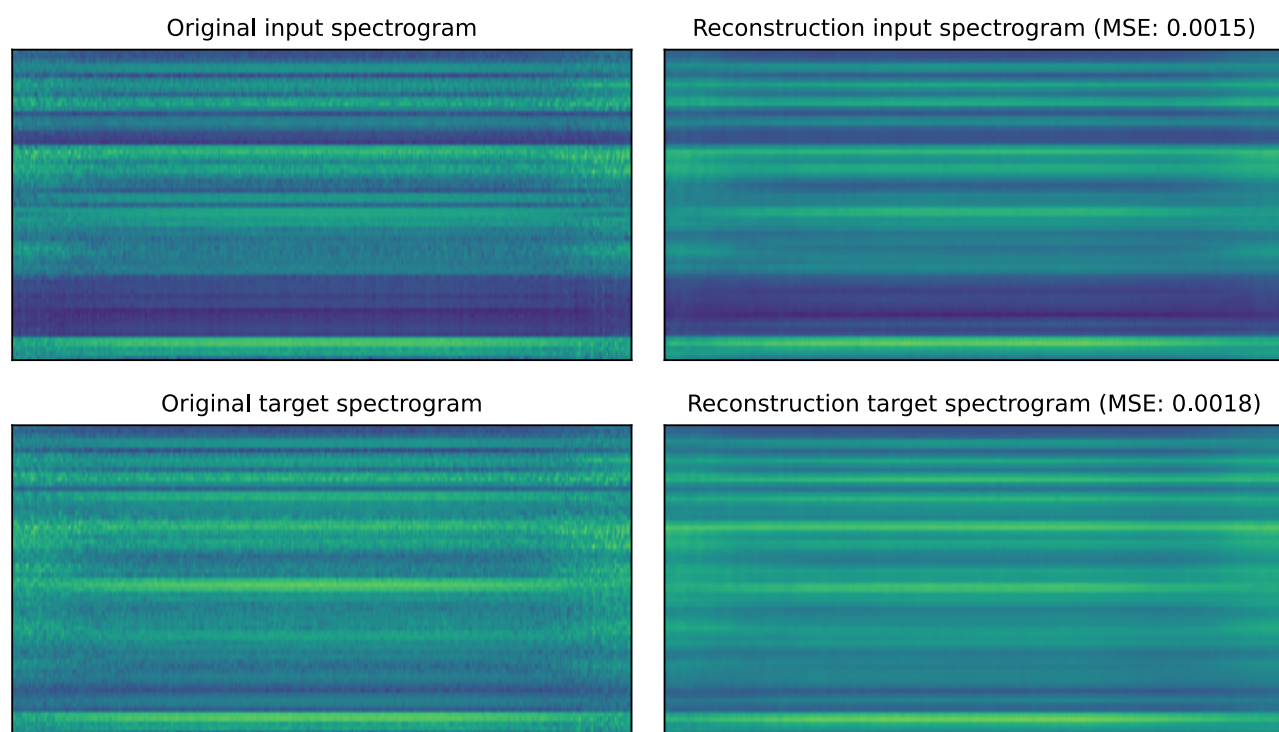


Figure 3: Reconstruction of Delay DAFX.

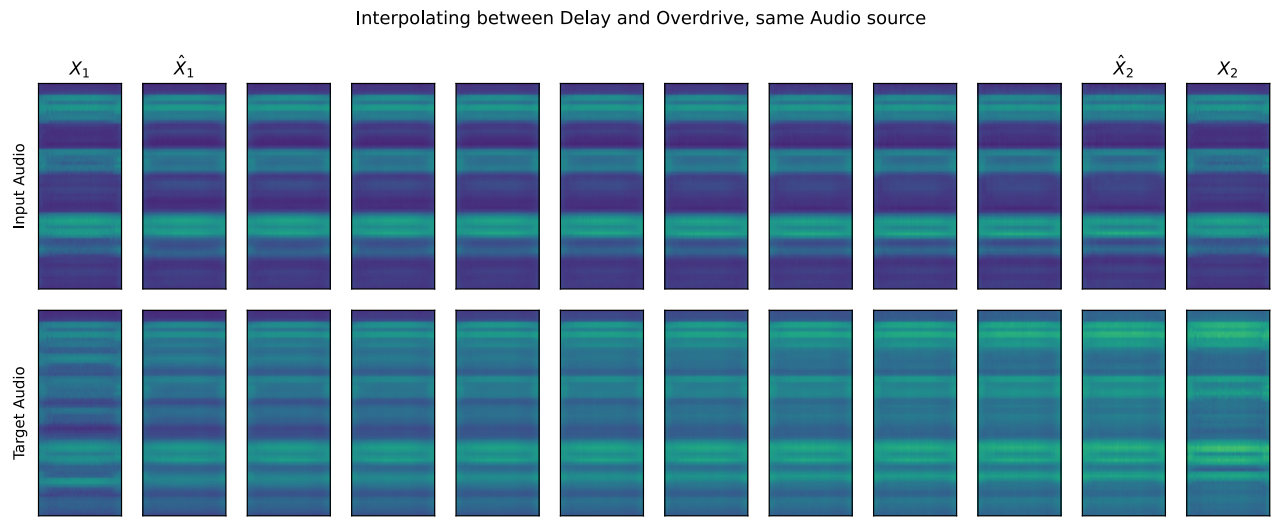


Figure 4: In this figure X_i is the original spectrogram and \hat{X}_i is the reconstruction from z_i (the latent embedding of X_i). The top row represents the 1st channel of the spectrogram (clean, unaffected audio) and the bottom row is the 2nd channel of the spectrogram (audio with DAFX applied).

The spectrograms between \hat{X}_1 and \hat{X}_2 are generated from linearly interpolating between z_1 and z_2 . Here, the audio source is the same for X_1 and X_2 , however, X_1 has the Delay effect applied, while X_2 has the Overdrive effect applied (hence, ideally the top row should not change throughout the interpolation).

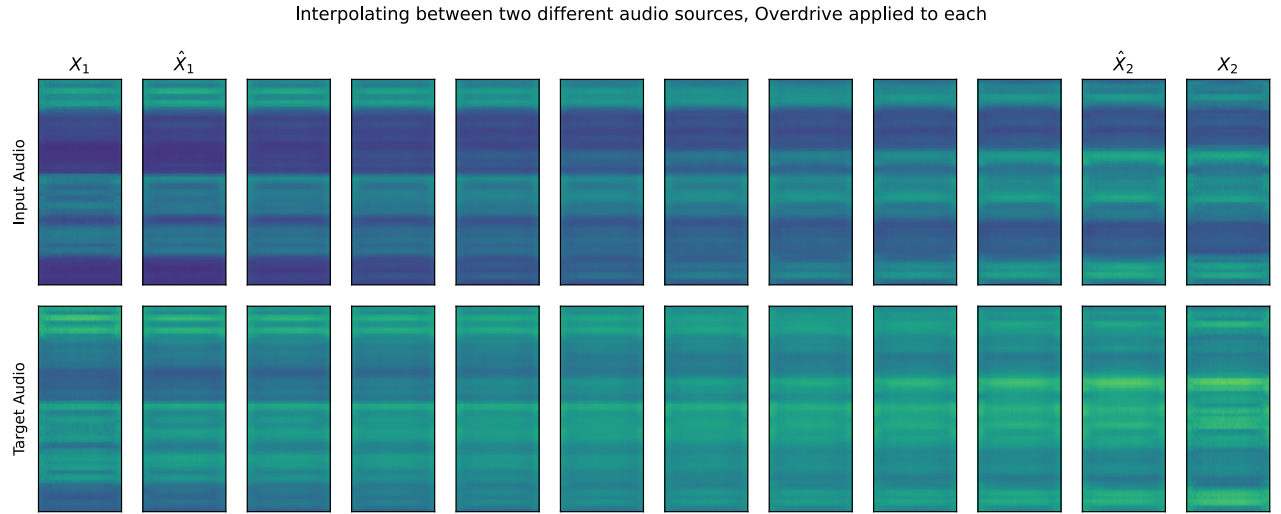


Figure 5: As in **Figure 4**, however in this image we use two different audio sources and apply a single effect to each (Overdrive).
