

Code snippet for normalising spectrograms:

```
...
X = torch.stft(
    signal.view(bs, -1),
    n_fft=n_fft,
    hop_length=hop_length,
    window=window,
    return_complex=True,
)

# Absolute value part
X_db = torch.pow(X.abs() + 1e-8, 0.3)
X_db_norm = X_db

# Normalise (0,1)
X_db_norm -= 0.3352797
X_db_norm /= 0.2745147
...
```

Interpolation of output_db setting for Overdrive DAFX




Figure 1: Interpolating between -20dB and +20dB with spectrogram normalisation applied. True spectrogram is shown on the top row, reconstruction is shown on the bottom row. In this case, the spectrograms are all very similar.




Figure 2: Cosine similarity of latent embeddings for the above interpolation.




Figure 3: Interpolating between -20dB and +20dB with spectrogram normalisation applied for a different audio sample. True spectrogram is shown on the top row, reconstruction is shown on the bottom row.

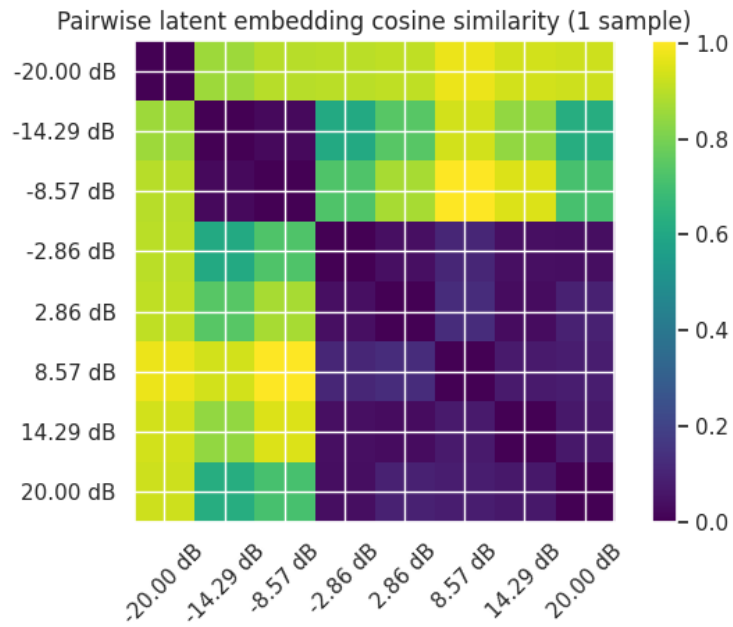


Figure 4: Cosine similarity of latent embeddings for the above interpolation.




Figure 5: Cosine similarity of latent embeddings averaged over 500 samples for each setting.
