

Week 10 Report

Kieran Grant - 2357351G

What I've done this week

- Tried changing out Spectrograms for Mel-Spectrograms.
 - Reduced frequency range to 20Hz-12kHz with 256 Mel-frequency bins (256x256 image).
- Also implemented a VQ-VAE to potentially replace the β -VAE.
 - Reconstructions were much better, but created higher dimensional latent embeddings.
 - Actual structure of latent space was not as good as previous Spectrogram β -VAE implementations (see Figures 1-3), so I've gone back to using those for now.
- Updated End-to-End model using the best Spectrogram-VAE so far (128D latent space).
 - Concatenating input and target embedding to create 256D input vector to the network.
 - Added Layer normalisation and initialisation so the model begins by predicting ~ 0.5 for all parameter settings at the start of training. Seems to have fixed issue of the model getting stuck from a few weeks ago.
 - Tested with static settings for a couple of DAFX.
 - Performed training runs with random settings for 4 DAFX. You can see plots and listen to the audio clips on W&B here: <https://api.wandb.ai/links/kieran-grant/kt0ixy4s>
- Kept working on latent controller interface.
 - Using Streamlit for the interactive plot; seems to work well!
 - User can click on datapoint in plot and the relevant audio is displayed.
 - Also created data generation pipeline to pre-compute embeddings (with various numbers of semi-supervised labels) as well as audio.

Questions

- Do you think it is worth starting the user evaluation with what I have, with the chance that the model *might* improve in the next couple of weeks?

Plan for next week

- Finish interface and test with expert user.
- Continue to try to improve model performance.
- Create scripts for other testing/MUSHRA.
- Continue writing report.

Current state of project

- Implementation is almost done, still looking to improve model performance however.
- Hoping to start evaluation as soon as possible.

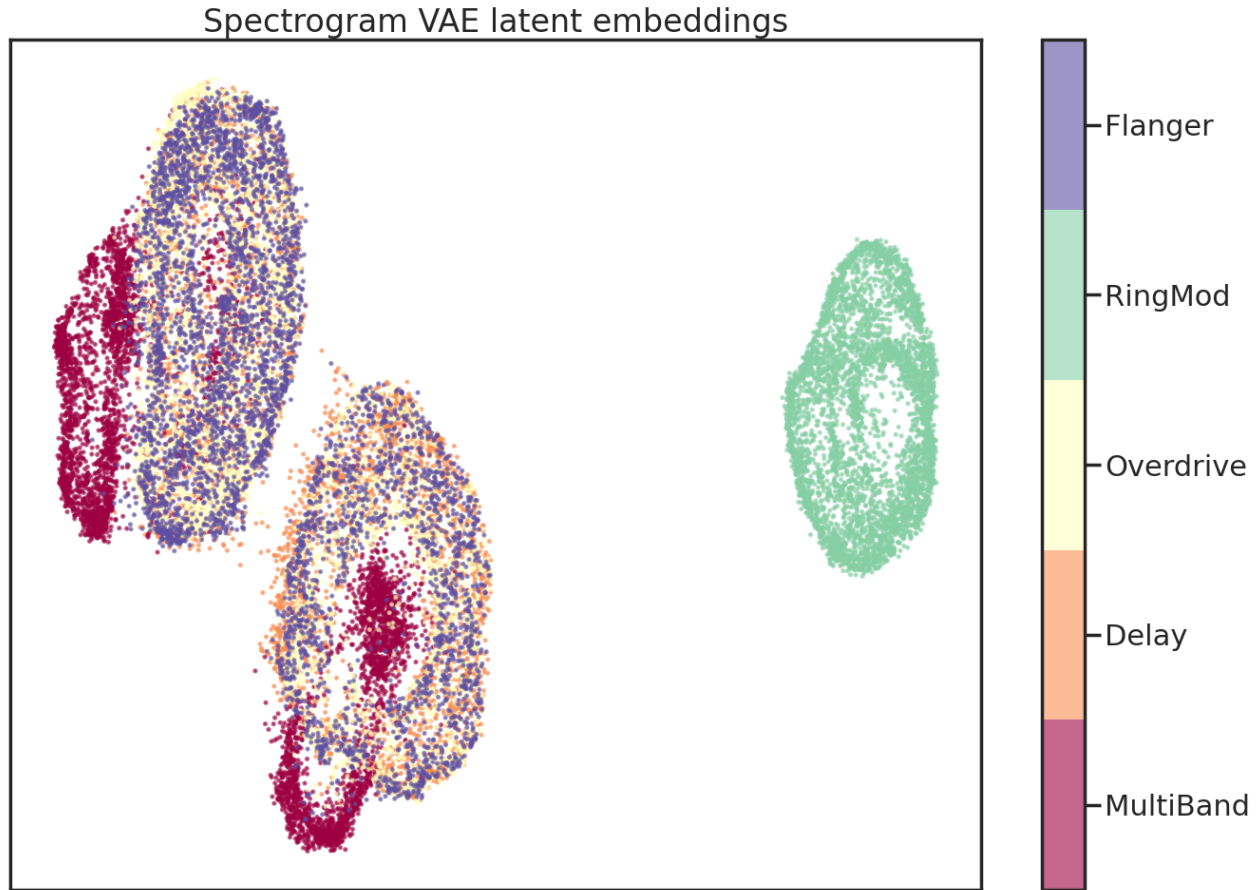


Figure 1: Latent embedding of 5 DAFX with VQ-VAE embeddings.

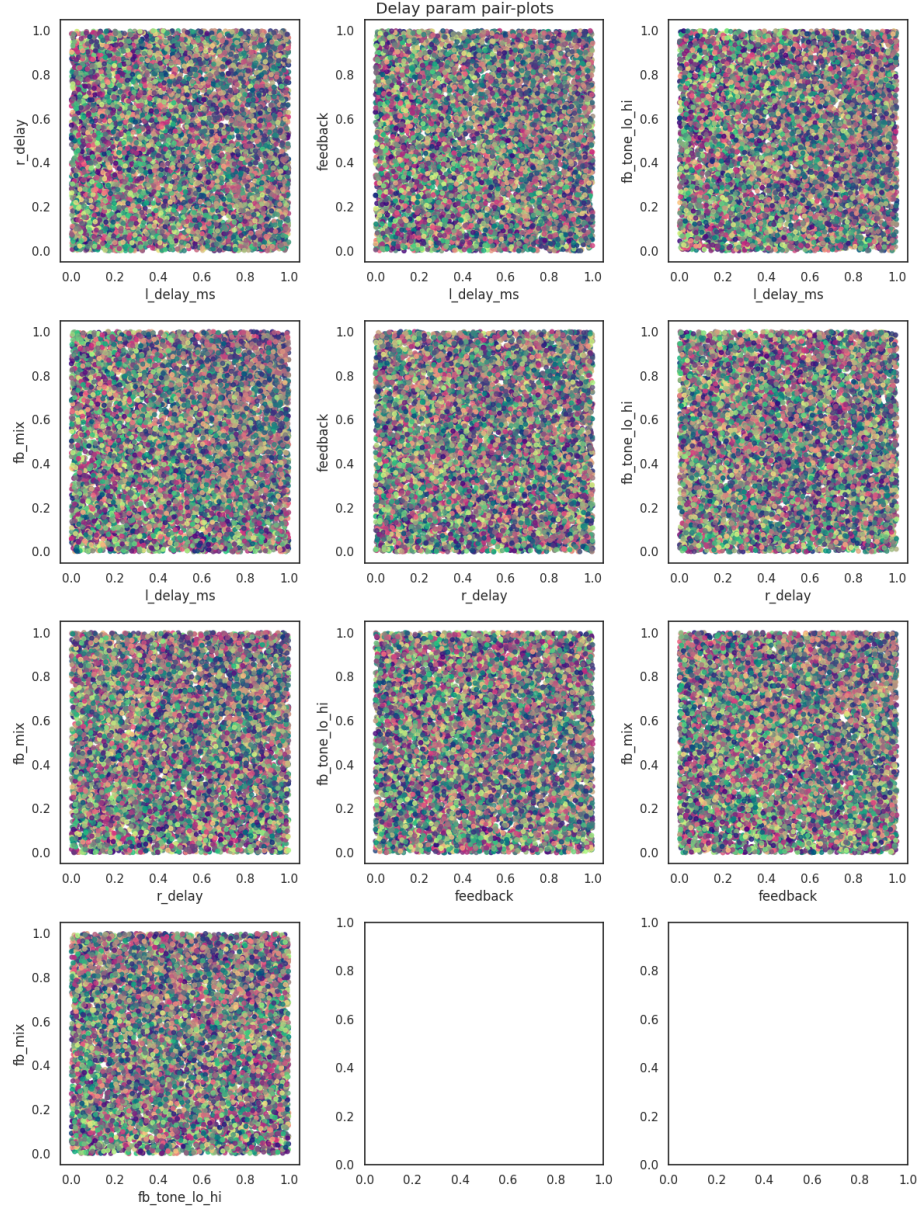


Figure 2: Domain colouring pairplots for Delay with VQ-VAE embeddings.




Figure 3: Domain colouring pairplots for Overdrive with VQ-VAE embeddings.
